

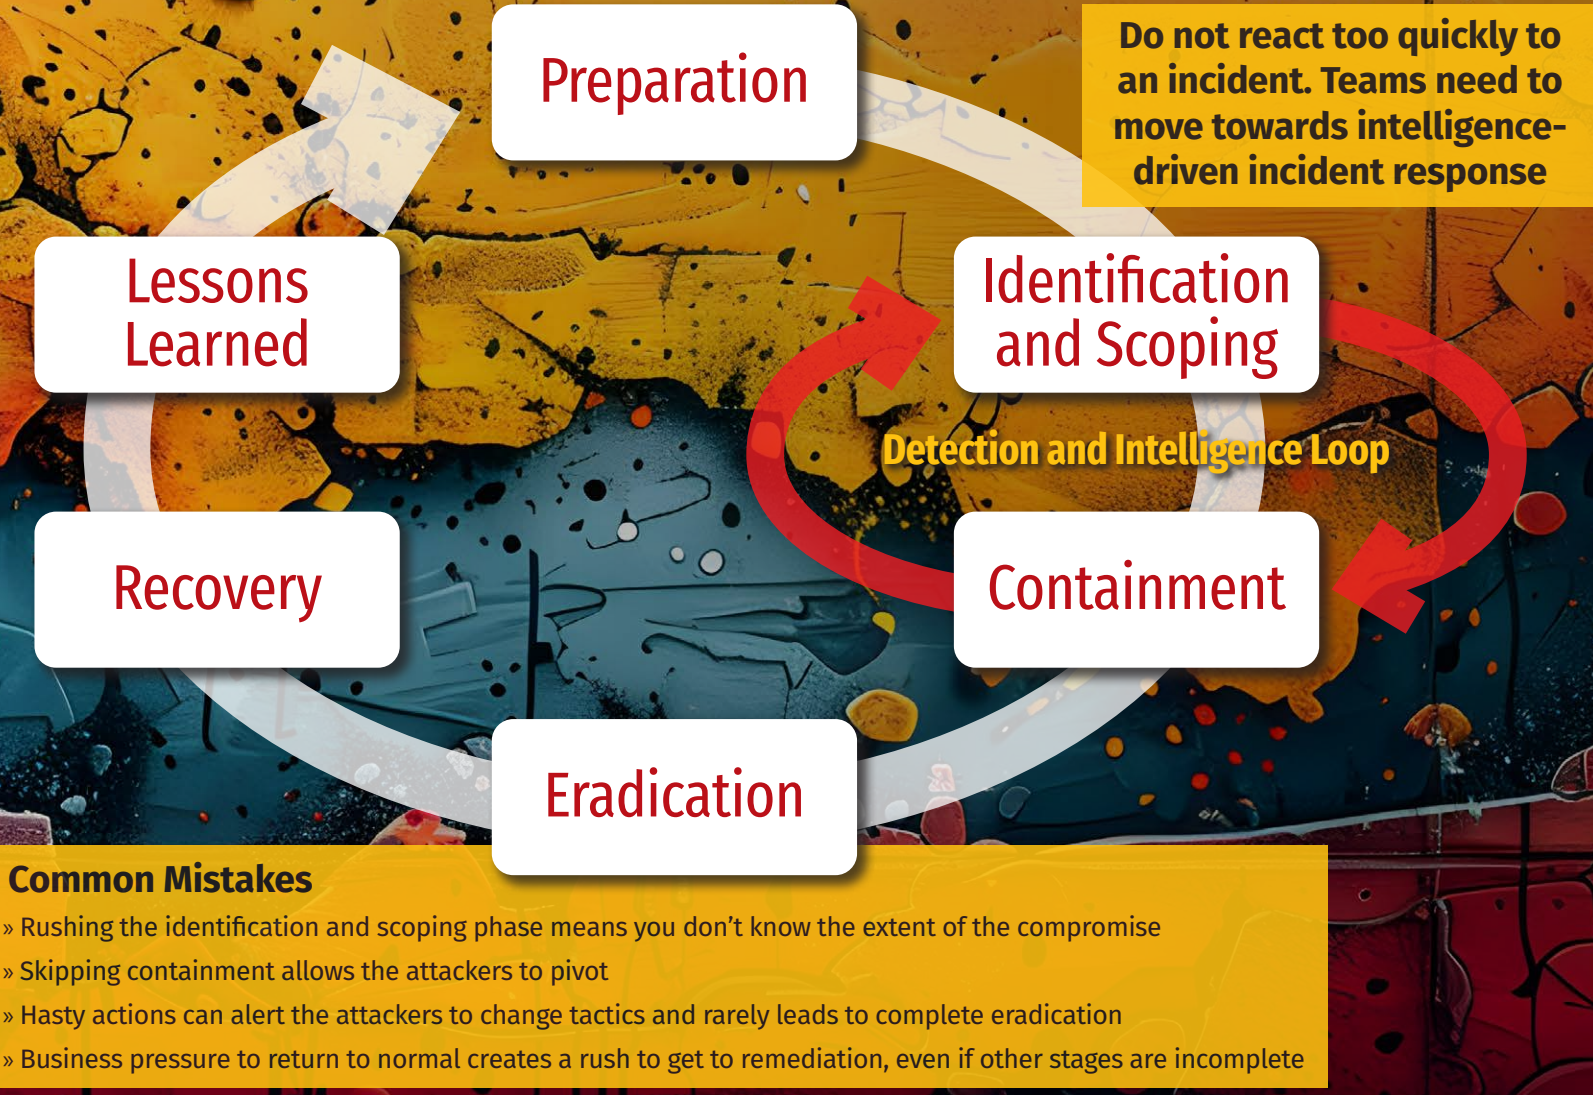
Linux Incident Response

digital-forensics.sans.org

Poster was created by Kathryn Hedley primarily based on the research and knowledge of Taz Wake in authoring FOR577.
©2025 SANS Institute. All Rights Reserved

DFIR_FOR577_0325

THE INCIDENT RESPONSE PROCESS



PACKAGE MANAGERS

- » **apt** = Advanced Package Tool
 - » Default on Ubuntu systems
 - » Tool for Debian systems
 - » Acts as a front end for Debian packages
 - » Manages dependencies
 - » Repository list (includes default repos): `/etc/apt/sources.list`
 - » e.g., added using the command `apt-add-repository`
- » **dpkg** = Debian Package
 - » Low-level tool for Debian systems
 - » Does not manage dependencies
 - » Typically used for the installation of standalone packages (.deb files)
 - » In general, .deb files are not digitally signed; trust is placed on the repository
 - » View the contents of a .deb file: `ar -tv [package]`
 - » List the contents of the data.tar* sub-file within a .deb file: `dpkg -c [package]`
 - » Extract the contents of a .deb file: `ar -xov [package]`
 - » Extract the contents of the data.tar* sub-file within a .deb file: `dpkg -x [package]`
 - » Display any metadata associated with a .deb file: `dpkg -f [package]`
 - » Note: The ar tool may not be installed by default; install the `binutils` package
- » **rpm** = RPM Package Manager
 - » Low-level tool for Red Hat-based distros
 - » Found on RHEL/CentOS
 - » Does not manage dependencies
 - » Effectively equivalent to dpkg on Debian
 - » Convert an .rpm file to a .cpio file: `rpm2cpio [RPM file] > [CPIO file]`
 - » View the contents of a .cpio file: `cpio -i -t -v < [CPIO file]`
 - » Extract an individual file from a .cpio file: `cpio -i --to-stdout [file] < [CPIO file]`
- » **yum** = Yellowdog Updated, Modified
 - » Default on Red Hat Enterprise Linux and CentOS
 - » High-level wrapper for the rpm tool
 - » Manages dependencies
 - » Python 2-based tool with some additional libraries
 - » Repository lists (includes default repos): `/etc/yum.repos.d/*repo`
- » **dnf** = Dandified YUM
 - » Default on Fedora
 - » Rewrite of yum
 - » Manages dependencies
 - » Uses libdnf framework, librepo, and libcomps
 - » Repository lists (includes default repos): `/etc/yum.repos.d/*repo`
- » **apk** = Alpine Package Keeper
 - » Lightweight package management tool used on Alpine Linux
 - » Designed for use in low-resource environments such as containers
 - » Manages dependencies
- » **snap**
 - » App Store and package manager developed by Canonical
 - » Snap applications run in a sandbox environment with limited access to the host
 - » Self-contained compressed filesystem using SquashFS
- » **synaptic**
 - » GUI for APT
 - » Used on Debian systems
 - » Uses deb and rpm packages
 - » Less common than snap
- » **zypper** = Zen/YaST Packages Patches Patterns Products
 - » Command-line interface for YaST (Yet another setup tool)
 - » Uses rpm-based packages
 - » Found on openSUSE and SUSE Linux Enterprise
 - » Repository lists (includes default repos): `/etc/zypp/repos.d/*`

Action	APT	DNF	ZYPPER	APK
Install PackageName	apt install PackageName	dnf install PackageName	zypper in PackageName	apk add PackageName
Update Package Database	apt update	dnf check-update	zypper ref	apk update
Remove PackageName	apt remove PackageName	dnf remove --nods PackageName	zypper rm -RU PackageName	apk del PackageName
Upgrade all packages	apt upgrade	dnf update	zypper up	apk upgrade
List available upgrades	apt list --upgradeable	dnf check-update	zypper lu	apk info (see notes)

Threat Hunting Package Managers

- » Review all installed applications
 - » Look for random or short (1-2 character) application names
 - » Look for application names that relate to known indicators
 - » Look for applications installed during any time period of interest
 - » Look for unusual repository use
- » Review individual packages
 - » Check internal file timestamps
 - » Check any signatures; they should be intact, valid, and expected
 - » Review any pre- or post-installation scripts to check for unusual behavior
- » Extract individual installer files and review them

USEFUL LINUX COMMANDS

- awk** Powerful data manipulation and pattern matching
- cat** Concatenate files or display text
- column** Format output with columns
 - t Table format output
 - S Specify an input separator (the default is whitespace)
 - o Specify an output separator (the default is two spaces)
 - j JSON output (requires column names to be specified)
 - N Specify the column names with a comma-separated list; this will be used for the table header
 - n Provide a name for the table (the default is "table")
- cut** Slice input to get the bits you want
 - d Use this as a delimiter to cut on
 - f Select this field or combination of fields
- date** Print or set the system time
 - d [STRING] Display the date referenced by STRING, not "now"
 - u Display the time in UTC
- file** Check the file's file signature
 - Z Look inside zip files/compressed archives
- grep** Search file contents for a pattern match
 - r Search recursively
 - i Case-insensitive search
 - o Print only the matched [non-empty] parts of a matching line
 - v Exclude matches
 - E Pattern match to an extended regular expression
 - a Force grep to treat all input as ASCII
 - A# Show # (number) lines after the matching line
 - B# Show # (number) lines before the matching line
 - C# Show # (number) lines before and after the matching line
- head** Show the first N lines of a file
 - n # Show # number of lines
 - c # Show # number of bytes
- hexdump** Display file contents in hex (or octal etc.)
 - C Canonical hex + ASCII display, which is similar output to xxd
 - n [number] Interpret only [number] bytes of input
 - s [number] Skip [number] of bytes from the start
- kill** Signal process
 - s Send a specific signal name or number
 - l Send the signal to all processes with a PID greater than 1
 - L Print a list of signal names
- ls** List directory contents
 - a Show hidden files
 - l Use long format and show more details
 - h Use human-readable numbers; 1024 = 1k, etc.
 - l Show inode numbers
- man** Display the manual for a command
 - k [keyword] Search the manual pages for entries that refer to the keyword
 - f [command] Display a concise description of the command
- md5sum filename** ... Generate an MD5 hash of filename
- sha1sum filename** ... Generate a SHA1 hash of filename
- sha256sum filename** ... Generate a SHA256 hash of filename
- sort** Takes input and sends to stdout in order
 - r Reverse the sort order
 - n Compare strings based on numeric value
 - b Ignore leading blanks
 - u Suppress all but one in each set where keys are equal
- tail** Show the last N lines of a file
 - n # Show # number of lines
 - c # Show # number of bytes
- tar** Manage archives
 - tf [file] List the contents of [file] archive
 - xf [file] Extract the contents of [file] archive
- uniq** Removes duplicate adjacent lines
 - d Only return the duplicate lines
 - c Count the number of occurrences
 - i Ignore the case
- wc** Count the number of bytes, lines, or words in a file
 - l Show the number of lines
 - c Show the number of bytes
 - w Show the number of words
 - m Show the number of characters
- xxd** Making hexdumps
 - s [number] Start [number] bytes into the file (seek)
 - l [number] Show [number] bytes in total (length)
 - r Convert hexdump back into a binary format (reverse)
- zcat** Read compressed files
 - d Decompress the file
 - r Recurse through directories
- zgrep** Search compressed data sets

Common Regex metacharacters:

- .
 - ^
 - \$
 - *
 - +
 - ?
 - []
 - |
 - ()
- Escape a metacharacter to search for the literal character using backslash (\)

INCIDENT RESPONSE AND THE SLEUTH KIT

Incident Response Task	TSK Command
Review an image to understand its structure.	<code>mmls [image]</code>
Identify the layout of a filesystem.	<code>fsstat -f [filesystem] -i [image type] [image]</code>
Extract all filesystem timestamps for a timeline.	<code>fls -r -m [image]</code>
Read the contents of a file with a given filename.	<code>fcats [file path] [image]</code>
Read the contents of a file at a given inode.	<code>icat [image] [inode]</code>
Review the metadata of a file.	<code>istat [image] [inode]</code>
Extract the contents of a data block.	<code>blkcat [image] [block number]</code>
Identify the filename for a given inode.	<code>ffind [image] [inode]</code>
Identify the inode for a given filename.	<code>ifind [image] -n [file path]</code>
Extract an individual disk partition.	<code>mmcat [image] [partition number] > [file]</code>
List unallocated inodes in Mactime format.	<code>ils -m [image]</code>
List all deleted regular files and their inodes.	<code>fls -fnd [image]</code>

GENERAL FILESYSTEM TERMINOLOGY

BLOCK	An area of storage space that is the smallest accessible unit on a filesystem.
BLOCK GROUPS	Groups of blocks containing filesystem-related information.
EXTENT	A group of consecutive blocks
SUPERBLOCK	The block that contains the filesystem metadata that defines other structures.
INODE	An "index node." This stores file metadata and its location on the drive.
INODE BITMAP	Records which inodes are in use within a group and the locations of the data block.

TIMESTAMPS

M = mtime	A = atime	C = ctime	B = btime = crtime
Modified: Data content change time	Accessed: Data last access time	Changed Metadata: Metadata change time	Birth: File creation time
» Time the file data was last modified	» Approximate time when the file data was last accessed for some reason, not necessarily by the user	» The time the inode entry was last updated	» Time file was created on this filesystem
	» Not reliable		» Not always recorded

General Linux Filesystem Time Rules

	File Creation	File Access	Content Change	File Copy ²	File Move/Rename	File Move (New Mount Point)	File Delete ³
M Modification	Set to time of creation	No change	Set to time of change	Set to time of copy	No change	No change	Set to time of deletion
A Access	Set to time of creation	Set to time of access ¹	Set to time of change ¹	Set to time of copy	No change	No change	No change
C Metadata Change	Set to time of creation	No change	Set to time of change	Set to time of copy	Set to time of move	Set to time of move	Set to time of deletion
B Creation	Set to time of creation	No change	No change	Set to time of copy	No change	Set to time of move	No change

¹ This is controlled by the atime settings when the filesystem is mounted.
² These timestamps are for the new file. The source file may have an atime change.
³ This varies by filesystem type. XFS does not set the mtime. EXT stores a 32-bit deletion timestamp. XFS does not.

FILESYSTEM KEY CHARACTERISTICS

Filesystem	Theoretical Maximum Single-File Size	Theoretical Maximum Filesystem Size	Timestamps	Journal	Notes
EXT3	2 TiB	16 TiB	32-bit, no crtime Subject to "2038 timestamp issue"	Yes	Older, found on infrastructure
EXT4	16 TiB	1 EiB	64-bit, crtime Subject to "2446 timestamp issue"	Yes	Default on Debian family
XFS	8 EiB	8 EiB	64-bit, crtime Subject to "2446 timestamp issue"	Yes	Default on RHEL family
BTRFS	16 EiB	16 EiB	64-bit, crtime Subject to "2446 timestamp issue"	CoW	Found on NAS devices and some Fedora/SUSE versions
ZFS	16 EiB	256 ZiB	True 64-bit, crtime Not subject to "2446 timestamp issue"	CoW	Sometimes found on HPC/Unix systems
UFS1	8 ZiB	8 ZiB	32-bit, no crtime	No	Mainly used on Unix systems
UFS2	8 ZiB	8 ZiB	True 64-bit, crtime	No	Mainly used on Unix systems

TIMELINING

Creating a Filesystem Timeline

- » Step 1: Create a bodyfile:
 - » From a simple source image: `fls -a -r -m [path prefix] [source image] > [bodyfile]`
 - » From a complex source image: `fls -a -r -m [path prefix] -f [image format] -on [partition offset] [source image] > [bodyfile]`
- » Step 2: Convert to human-readable format:
 - » Convert to CSV: `mactime -d -z UTC -b [bodyfile] [time slice] > [timeline.csv]`
- » Step 3: Analyze

Creating a Super Timeline

- » Option 1: psteval
 - » `psteval.py --source [source image] -o [output format] -w [output file]`
- » Option 2: log2timeline and psort
 - » Extract events from image into storage file: `log2timeline.py --storage-file [storage file] [source image]`
 - » Create a human-readable timeline from the storage file: `psort.py -w [output file] [storage file] [event filter]`

ACQUIRING DATA

Common Disk Evidence Collection Tools

- » **dd**
 - » Installed by default in most distros
 - » Usage: `dd if=[input file] of=[output file] bs=[blocksize] conv=sync,noerror`
- » **dcfldd**
 - » Not installed by default
 - » Forked version of dd
 - » Usage: `dcfldd if=[input file] of=[output file] hash=[hash format] hashlog=[hashlog file] conv=sync,noerror`
- » **dc3dd**
 - » Not installed by default
 - » Patched, more verbose version of dd
 - » Usage: `dc3dd if=[input file] of=[output file] hash=[hash format] hlog=[hashlog file] log=[diagnostic log file]`
- » **ewf4quire**
 - » Not installed by default
 - » Part of libewf
 - » Writes E01 format files
 - » Usage: `ewf4quire [device]`

Common Manual Remote Collection Tools

- » **dd over ssh**:
 - » On the source system: `dd if=[input file] | ssh [username]@[ip address] "dd of=[output file] bs=[blocksize] conv=sync,noerror"`
- » **dd image sent to a Google Cloud Platform storage bucket**:
 - » On the source system: `dd if=[input file] bs=[blocksize] | gcp | gsutil cp -gs://[bucketname]/[objectname]`
- » **dd over netcat**:
 - » On the source system: `dd if=[input file] bs=[blocksize] | bzip2 -c | nc [ip address] [port]`
 - » On the destination system: `nc -nlp [port] | bzip2 -d | dd bs=[blocksize] of=[output file]`

Acquiring RAM

- » **Acquire Volatile Memory for Linux (AVML)**:
 - » Open-source tool written in Rust, available from: <https://github.com/microsoft/avml>
 - » Does not require on-target compilation of the binary
 - » Includes support for various cloud platforms, such as AWS, Microsoft Azure, and Google Cloud, to store memory images
 - » Supports compression of memory images using the `--compress` option
 - » Usage: `avml [output file]`
- » **Loadable Kernel Module (LIME)**:
 - » Requires on-target compilation of the binary, or can be compiled externally using the same distro version and kernel release and build process as the target, but this can be complicated
 - » This project is no longer maintained
- » **Direct access to the system's physical memory in /dev/mem**:
 - » Direct access to physical memory on a running system can potentially cause system crashes or other damage
 - » When power is lost, this ceases to exist and is not recovered in standard disk images
- » **Access to the virtual filesystem representing RAM at /proc**:
 - » The proc filesystem contains a hierarchy of special files that represent running processes and various system resources
 - » `cmdline`: process command line and any arguments—use `cat` to view
 - » `comm`: process command name; similar to `cmdline` but doesn't include the name of any application that called the process—use `cat` to view
 - » `cwd`: current working directory that the process is running in—use `ls` to view
 - » `environ`: process environment assigned to the running application—use `cat` to view
 - » `exe`: symbolic link to the executable running the process—use `ls` to view
 - » `fd`: folder containing details on where the file descriptors are pointing—use `ls -al` to view
 - » `map_files`: folder containing links to shared objects and memory-mapped files used by the application—use `ls -al` to view
 - » `maps`: file that contains similar information to the `map_files` folder. Use `cat` to view
 - » `mountinfo`: file that holds information about any active mount points. Use `cat` to view
 - » `net`: directory containing networking information for the process—this includes TCP/UDP ports, route tables, and network statistics—use `ls` to view
 - » `root`: link to the root folder for this application—this can be useful for identifying applications running in unusual environments, such as chroot—use `ls -al` to view
 - » `sched`: file containing the CPU scheduling parameters as they apply to the PID being reviewed—use `cat` to view
 - » `stack`: file that describes the current kernel stack of the process being reviewed—use `cat` to view
 - » `status`: file containing the current status of a process—use `cat` to view
 - running (active)
 - sleeping (waiting for an event or waiting for a resource to become available)
 - zombie (this is a process that has completed execution but not yet fully exited the process list)
 - » Information is provided in real time and is dynamically updated by the kernel
 - » Can use filesystem tools to analyze
 - » When power is lost, this ceases to exist and is not recovered in standard disk images
- » **/proc/kcore** is a snapshot of kernel and physical memory

GENERAL LINUX FOLDER HIERARCHY

The diagram illustrates the general Linux folder hierarchy, starting from the filesystem root (/). The root branches into various top-level directories: boot, proc,/sbin,bin,lib,tmp,etc,root,home,usr,opt,var,cdrom,mnt,media,run,dev. A legend on the left defines the color-coding: Virtual FS (blue), Non-persistent (black), Libraries (orange), Binaries (yellow), User files (green), Variable files (red), and Runtime (yellow). Dashed lines indicate sub-structures: /bin and /sbin link to /usr/bin and /usr/sbin; /lib links to /usr/lib and /usr/share; /usr branches into include, local, share, and src; /var branches into cache, log, spool, run, lib, opt, and temp; /run links to /var/run; and /dev links to /dev/shm.

```
graph TD
    Root[" / filesystem root "]
    Root --> boot
    Root --> proc
    Root --> sbin
    Root --> bin
    Root --> lib
    Root --> tmp
    Root --> etc
    Root --> root_dir["root"]
    Root --> home
    Root --> usr
    Root --> opt
    Root --> var
    Root --> cdrom
    Root --> mnt
    Root --> media
    Root --> run
    Root --> dev

    subgraph "usr structure"
        usr --> usr_bin["bin"]
        usr --> usr_sbin["sbin"]
        usr --> usr_lib["lib"]
        usr --> usr_include["include"]
        usr --> usr_local["local"]
        usr --> usr_share["share"]
        usr --> usr_src["src"]
    end

    subgraph "var structure"
        var --> var_cache["cache"]
        var --> var_log["log"]
        var --> var_spool["spool"]
        var --> var_run["run"]
        var --> var_lib["lib"]
        var --> var_opt["opt"]
        var --> var_temp["temp"]
    end

    run -.-> var_run
    dev -.-> dev_shm["shm"]

    style boot fill:#ccc
    style proc fill:#000,color:#fff
    style sbin fill:#00aaff,color:#fff
    style bin fill:#ffff00
    style lib fill:#ffa500
    style tmp fill:#ccc
    style etc fill:#ccc
    style root_dir fill:#ccc
    style home fill:#ccc
    style usr fill:#00aaff,color:#fff
    style opt fill:#ccc
    style var fill:#ff0000,color:#fff
    style cdrom fill:#ccc
    style mnt fill:#ccc
    style media fill:#ccc
    style run fill:#ffff00
    style dev fill:#ccc
```

KEY

- Virtual FS
- Non-persistent
- Libraries
- Binaries
- User files
- Variable files
- Runtime

Some important locations for DFIR:

- > /bin and /sbin – binary executables
- > /etc – configuration files
- > /dev – device files
- > /home – user home directories
- > /lib* – libraries and system files
- > /mnt and /media – mounted filesystems
- > /opt – optional third-party software packages
- > /root – root user's home directory
- > /tmp – temporary files
- > /usr – user-related data, including executables and libraries
- > /var – log files, and temporary or persistent data for various services and applications

Types of image file	Raw Image Files	VMDK and VHDX files
<ul style="list-style-type: none"> ▶ Raw (often referred to as DD) <ul style="list-style-type: none"> ▶ Often with an extension: <code>dd, .raw, .001, .img</code> ▶ Easiest to mount; only requires the built-in <code>mount</code> command ▶ .E01 <ul style="list-style-type: none"> ▶ EnCase Expert Witness Format (EWF) ▶ Industry standard in forensic evidence image formats ▶ Files normally have an <code>.E01</code> extension, but this can vary if the file is split ▶ Files need to be converted to raw format first—this can be achieved using <code>ewfmount</code>, which is not installed by default 	<ul style="list-style-type: none"> ▶ Mounting a raw image file with one partition: <ul style="list-style-type: none"> ▶ <code>mkdir [dir] to use as mount point</code> ▶ <code>mount -o [raw image file] [mount point dir]</code> <ul style="list-style-type: none"> ▶ The <code>-o</code> option mounts the image in read-only mode. ▶ Mounting a raw image file with multiple partitions: <ul style="list-style-type: none"> ▶ <code>fsck -l [raw image file] OR mmls [raw image file]</code> <ul style="list-style-type: none"> ▶ Identify the partition you want to mount and multiply the partition start location by the sector size to get the offset. ▶ <code>mount -o ro,offset=[offset] [raw image file] [mount point dir]</code> ▶ Unmounting a raw image file: <ul style="list-style-type: none"> ▶ <code>umount [mount point]</code> ▶ <code>rmdir [mount point dir]</code> ▶ Mounting an LVM/LVM2 partition: <ul style="list-style-type: none"> ▶ <code>fsck -l [raw image file] OR mmls [raw image file]</code> <ul style="list-style-type: none"> ▶ Identify the partition you want to mount and multiply the partition start location by the sector size to get the offset. ▶ <code>losetup -rf -o [offset] [raw image file]</code> ▶ <code>losetup -a grep [raw image file]</code> 	<ul style="list-style-type: none"> ▶ Mounting a VMDK file: <ul style="list-style-type: none"> ▶ <code>mkdir [dir] to use as mount point</code> ▶ <code>guestmount -a [VMDK file] -i -ro -o allow_other [mount point]</code> <ul style="list-style-type: none"> ▶ <code>-a</code> Specifies the source image to be mounted ▶ <code>-i</code> Uses inspector to attempt to detect the operating system and mount points automatically ▶ <code>-ro</code> Mount the image as read-only. ▶ <code>-o Options</code> ▶ <code>-allow, other:</code> Allow other accounts to see the mounted filesystem ▶ Mounting a VHDX file: <ul style="list-style-type: none"> ▶ <code>mkdir [dir] to use as mount point</code> ▶ <code>guestmount -a [VHDX file] -m [mount point in VHDX] -ro -o allow_other [mount point]</code> <ul style="list-style-type: none"> ▶ <code>-a</code> Specifies the source image to be mounted ▶ <code>-i</code> Uses inspector to attempt to detect the operating system and mount points automatically ▶ <code>-m</code> Specifies the mount point within the VHDX file to mount ▶ <code>-ro</code> Mount the image as read-only.

- Mounting an EWF image file with one partition:
 - » `mkdfr [dir to use as mount point for ewfmount command]`
 - » `ychange -a y [VG Name]`
 - You can, generally, ignore warnings about old PV headers.
 - » `lvs can | grep [VG Name]`
 - Validate the Volume Group is ACTIVE and identify root device to mount.
- Mounting an EWF image file with multiple partitions:
 - » `mkdfr [dir to use as mount point for ewfmount command]`
 - » `ewfmount [E01 image file] [mount point for raw image]`
 - » `mkdfr [dir to use as mount point for mount command]`
 - » `fdisk -l [raw image file] OR mmls [E01 image file]`
 - Identify the partition you want to mount and multiply the partition start location by the sector size to get the offset
 - » `mount -o ro,norecovery,offset={offset} [raw image file] [mount point]`
- Unmounting an EWF image file:
 - » `umount [raw mount point]`
 - » `rmrdr [raw mount point dir]`
 - » `umount [EWF mount point]`
 - » `rmrdr [EWF mount point dir]`
- Identifying the Volume Group Name (VG Name):
 - » `ychange -a y [VG Name]`
 - » `lvs can | grep [VG Name]`
 - Validate the Volume Group is ACTIVE and identify root device to mount.
 - » `mkdfr [dir to use as mount point]`
 - » `mount -o ro, norecovery, noexec {device to mount} [mount point dir]`
- Mounting an LVM/LVM2 partition:
 - » `umount [mount point]`
 - » `ychange -a n [VG Name]`
 - » `losetup -d [loopback address]`
 - » `rmrdr [mount point dir]`
- Unmounting a VMDK or VHDX file:
 - » `guestmount [mountpoint]`

MODIFIED USER ACCOUNTS	STRANGE EVENTS IN SHELL HISTORY	TEXT EDITOR HISTORY FILES	STRANGE NETWORKING	COMMON PERSISTENCE	ALTERED FILES	SUSPICIOUS LOG DATA
<h2>Account Modification</h2> <p>➤ Signs of Attack:</p> <ul style="list-style-type: none"> ➤ Brute-force attacks ➤ Creating new accounts ➤ Adding accounts to groups to elevate permissions ➤ Adding SSH keys to the authorized_keys file ➤ Giving themselves sudo rights <p>➤ Locations:</p> <ul style="list-style-type: none"> ➤ Authentication logs (auth.log, secure, utmp, wtmp, btmp, etc.). Check for large numbers of failed logins ➤ /etc/passwd: Check the user accounts, look for file last modification times, and check for users who have long shell firsts ➤ Readable by all users ➤ File structure: <code>{Username}{Password info (x = stored in shadow)}:UID (UID){Group ID (GID)}:User information (comments){Home folder}{Login shell}</code> ➤ /etc/shadow: Same as the passwd file, look for any unexpected accounts and check the time the file was last modified. Check for unusual entries in the password field. ➤ Only readable by root user ➤ File structure: <code>{Username}{Password info}{Last password change}{Minimum password age}{Maximum password age}{Password warning period}{Password inactivity period}{Account expiration date}{Reserved field}</code> <ul style="list-style-type: none"> • The {Password info} field has the structure: <code>\$1\$hash\$algorithm code\$Salt\$hash</code> • If the {Password info} field contains an asterisk (*) or exclamation marks, the account is locked and cannot be used for password-based login. • The last three fields are rarely used. ➤ /etc/group: Check what groups exist and which accounts are in privileged groups (start with wheel, sudo, and adm). Check the last modification time. ➤ /etc/sudoers: Validate the last modification time. Check for users with excessive accounts, especially users who have ANY_Any set with or without a password. The sudoers file is often overwritten by the system as part of updates/patches. ➤ /etc/sudoers.d/*: As for the sudoers file, attackers prefer this location because it survives patching/system updates ➤ /home/{username}/ssh/ and /root/.ssh/ <ul style="list-style-type: none"> ➤ The known_hosts file stores a record of public keys for systems the user account has connected to. ➤ The authorized_keys file contains public keys that are allowed to authenticate on this system. 	<h2>Shell and Text Editor History</h2> <p>➤ Signs of Attack:</p> <ul style="list-style-type: none"> ➤ Use of sudo/su ➤ Cleartext passwords ➤ Unusual command lines ➤ Locations: <ul style="list-style-type: none"> ➤ bash_history or zsh_history (shell history): Look for any unusual commands that have been run ➤ The location of this file is stored in the HISTFILE variable ➤ Only written to disk on shell exit (if exit recognized by the OS) ➤ Order of commands is not reliable ➤ Trivial for an attacker to edit or remove ➤ May be truncated based on the current HISTFILESIZE setting, with oldest entries removed first ➤ lesshist: Look for unusual searches that were conducted ➤ viminfo: Look for any unusual use of vi/vim ➤ mysql_history: Look for any unusual mysql activity ➤ Any other history files such as python_history, gdb_history, or local/share/nano/search_history 	<h2>Networking</h2> <p>➤ Signs of Attack:</p> <ul style="list-style-type: none"> ➤ Unusual ports being used ➤ Long-lasting connections ➤ Unexpected processes using the network <p>➤ Locations:</p> <ul style="list-style-type: none"> ➤ /etc/hosts: Check that hostnames point to legitimate targets and haven't been modified by an attacker ➤ /etc/resolv.conf and /etc/system/resolved.conf: Check that the nameserver settings are valid 	<h2>Altered files</h2> <p>➤ Signs of Attack:</p> <ul style="list-style-type: none"> ➤ Adding users ➤ Deleting history/anti-forensics ➤ Hiding files ➤ Staging data ➤ Leaving backdoors <p>➤ Locations:</p> <ul style="list-style-type: none"> ➤ Look for large files, especially archives. This can often be a sign that attackers are staging data for exfiltration. ➤ /dev – Look for regular files in here; it should only hold devices or links ➤ Search for files with a modification time stamp that fits within the incident window. Remember that timestamps can be manipulated. 	<h2>Log Files</h2> <p>➤ Signs of Attack:</p> <ul style="list-style-type: none"> ➤ Failed attempts to use sudo ➤ Root account changing key system settings ➤ New accounts/cron jobs being created ➤ Unusual log entries during the incident window <ul style="list-style-type: none"> ➤ In general, the most effective way to work through logs is to follow the line of attack. ➤ Log files are typically very noisy. ➤ Maliciously modified log entries ➤ Deleted log entries or entire log files ➤ Attack to logging service causing log corruption <p>➤ Locations:</p> <ul style="list-style-type: none"> ➤ /var/log/ ➤ /var/run/ (symlink) ➤ Log management tool (logrotate) configuration: <ul style="list-style-type: none"> ➤ Configuration file: /etc/logrotate.conf ➤ Individual configurations: /etc/logrotate.d/* ➤ Scheduled task: /etc/cron.daily/logrotate ➤ Systemd timer: /usr/lib/systemd/system/logrotate* ➤ State file (last rotation time): <ul style="list-style-type: none"> • Ubuntu: /var/lib/logrotate/status • RHEL: /var/lib/logrotate/status 		

ROOTKITING

Types of Rootkit

- › **User-mode Rootkits:**
 - › Operate in user space, where regular applications run
 - › Typically hook application and system binaries to intercept and alter system calls
- › **Kernel-mode Rootkits:**
 - › Most popular rootkits
 - › More powerful and stealthy than user-mode rootkits
 - › Operate within kernel space, giving direct access to hardware and kernel subsystems
 - › They often modify kernel data structures or load as malicious kernel modules
 - › Techniques include hiding their presence by tampering with kernel data structures such as linked lists and process tables
 - › The most commonly found attack is the abuse of Loadable Kernel Modules (LKMs)
 - › Review Loadable Kernel Modules in the live kernel with the lsmod command. This command draws its information from the `/proc/modules` container.
- › **Bootloader Rootkits:**
 - › Infect the bootloader, such as GRUB (GNU GRand Unified Bootloader), which allows launch on boot
 - › They can initiate before the Linux kernel loads, making detection very difficult and allowing them to manipulate the boot process
- › **Hardware/Firmware Rootkits:**
 - › Target the firmware on devices
 - › These rootkits can persist through operating system reinstallations and have control from the earliest stages of the boot process
 - › Particularly challenging to detect and remove because they operate below the operating system level
- › **Memory Rootkits:**
 - › Reside solely in the system's RAM
 - › Typically inject code into the running kernel and manipulate its execution while the system is running
 - › Harder to detect but will not persist after a reboot unless reinjected
- › **Library Rootkits:**
 - › Target system libraries, replacing them or modifying their functions
 - › Can intercept and modify inputs and outputs of library functions, altering the behavior of applications using the libraries without altering the applications directly

Checking for rootkits

- › **Check for signs of known Linux rootkits on a live system:**
 - › `#!/usr/sbin/chkrootkit [options][test]`
 - › `rkhunter [options]`
- › **Review Loadable Kernel Modules:**
 - › Does the file name look legitimate?—malicious LKMs are often single-character names or contain types from legitimate modules
 - › Check the kernel log file (for example, `dmesg`) for anything unusual like tainting
 - › Where possible, compare known-good lists of loaded modules to the evidence
- › **Review running processes for library attacks:**
 - › Check and validate the paths to the files are legitimate in the environment variables for the running process—this information is stored in `/proc/$pid/environ`
 - › Use `ldd` to review an executable on disk and identify which shared libraries it calls compared to the libraries in the running process—this information is stored in `/proc/$pid/maps`
- › **Review system calls.**
 - › Review the global system call table for anything unusual
 - › Use `strace -p <PID>` to connect a debugger to Process ID <PID> and validate what system calls each process is making
 - › Use `strace <command>` to run a new command to check its system calls
- › **Check the audit logs, if configured**
- › **Scan the environment with Yara rules, to look for known indicators**

FOR577 teaches the skills needed to identify, analyze, and respond to attacks on Linux platforms and how to use threat hunting techniques to find the stealthy attackers who can bypass existing controls. The course addresses today's incidents by teaching the hands-on incident response and threat hunting tactics and techniques that elite responders and hunters are successfully using to combat real-world breach cases.

www.sans.org/for577

<h3>System Logs – Kernel Messages (dmesg)</h3> <p>➤ Description</p> <p>The Kernel Messages log contains information related to the boot process of the system, generated by the kernel and various drivers. The name of dmesg comes from the diagnostic messages. Entries in this log provide information about hardware and device initialization.</p> <p>➤ Location</p> <ul style="list-style-type: none"> ➤ <code>/var/log/dmesg</code> ➤ Also, on Ubuntu 22.04 and older: <code>/var/log/kern.log</code> <p>➤ Interpretation</p> <ul style="list-style-type: none"> ➤ The dmesg file is written at the end of the system startup cycle ➤ The log is stored in a buffer in the kernel and can be viewed using the dmesg command on a live system ➤ RHEL8/CentOS8 (Fedora 28), and Ubuntu 24.04 or newer, no longer include this service by default; dmesg events are written to the system journal and can still be viewed with the dmesg command, but there is no <code>/var/log/dmesg</code> file ➤ Includes entries relating to USB devices or drive activity 	<h3>System Logs – The Journal</h3> <p>➤ Description</p> <p>The Journal provides an audit trail of activity on the system by logging various types of events, including system startup and shutdown, kernel messages, and system service and application events. It stores metadata, such as the hostname, system architecture, and boot ID, allowing responders to correlate activity across systems and time.</p> <p>➤ Location</p> <ul style="list-style-type: none"> ➤ <code>/var/log/journal/[UUID]/*log</code> <p>➤ Interpretation</p> <ul style="list-style-type: none"> ➤ Logs are stored in a binary format ➤ In live response, the journal can be analyzed using the journalctl tool or programmatically via the systemd-journal API ➤ There are a lot of command line arguments and options available, some of the more useful ones for DFIR are: <ul style="list-style-type: none"> ➤ <code>-S or --since=</code> Show entries on or newer than the date specified (yyyy-mm-dd hh:mm:ss) ➤ <code>-l or --until=</code> Show entries on or older than the date specified ➤ <code>-b or --boot=</code> Show messages from a specific boot session ➤ <code>-u</code> Show messages for the specified systemd unit ➤ <code>-t</code> Show messages for the specified syslog identifier ➤ <code>-.COMM=</code> Match the script name ➤ <code>-EXEC=</code> Match the executable name ➤ <code>-g</code> Pattern match (grep) for the specified regular expression ➤ <code>-k or --dmesg</code> Show only kernel messages ➤ The journalctl command can also analyze offline journal 	<h3>Application Logs – vsftpd</h3> <p>➤ Description</p> <p>vsftpd is one of the more commonly found enterprise FTP clients on Linux systems, which unlike most other FTP packages, provides a secure method for transferring files over a network. It is configurable and provides robust logging capabilities.</p> <p>➤ Location</p> <p>Typically, logging is defined by the configuration file at: <code>/etc/vsftpd.conf</code> or <code>/etc/vsftpd/vsftpd.conf</code></p> <p>Log location:</p> <ul style="list-style-type: none"> ➤ <code>/var/log/vsftpd.log</code> <p>➤ Interpretation</p> <ul style="list-style-type: none"> ➤ The location of the configuration file on a live system can be established with the systemctl status vsftpd command ➤ By default, vsftpd logs events to the syslog facility ➤ By default only events of notice level or higher are logged
<h3>System Logs – Boot History</h3> <p>➤ Description</p> <p>Linux distros maintain a record of scripts run and events that fire during the boot process, in log files. This therefore logs the last time the system was booted.</p> <p>➤ Location</p>		<h3>Application Logs – Samba (SMB)</h3> <p>➤ Description</p> <p>Samba uses the SMB protocol to allow Linux/Unix servers to share files, printers, and other resources with Windows-based clients. Samba logs important events in its operation, including user authentication attempts and server status changes.</p> <p>➤ Location</p>

- **Description**
Account creation, user logins from external services and privilege use such as using sudo are recorded in an authentication log. The log location differs depending on the distro.
- **Location**
 - » Ubuntu: `/var/log/auth.log`
 - » RHEL: `/var/log/secure`
- **Interpretation**
 - » Plaintext log file
- **Description**
The most basic host firewall in Linux is iptables. It provides a way to control incoming and outgoing network traffic based on a set of rules. If the LOG option is used, information on packets matching the rule will be logged to the kernel message log.
- **Location**
 - » System Logs – Kernel Messages (dmesg)
- **Interpretation**
 - » On a live system, iptables logs in the kernel message log can be viewed using one of the commands:
 - » `dmesg`
 - » `journalctl -t iptables`
 - » See the Kernel Messages (dmesg) section for more information on these logs

<ul style="list-style-type: none"> » <code>/var/log/[Application]/ssl_access_log</code> » <code>/var/log/[Application]/ssl_request_log</code> » <code>/var/log/[Application]/ssl_error_log</code> 	<p>» Description</p> <p>UFW is a user-friendly front-end for iptables and a default firewall application in many Linux distributions, including Ubuntu. UFW is a default deny firewall, meaning it blocks all incoming traffic by default. Rules are then added to allow specific types of traffic.</p> <p>» Location</p> <p>» <code>/var/log/ufw.log</code></p> <p>» Interpretation</p> <p>» Logging is not turned on by default and must be enabled by the root account using ufw logging on</p> <p>» Log data generated by UFW (and iptables) can be noisy</p> <p>» Log files are plaintext, so easily readable with a text editor or using the cat command</p>
<p>» Interpretation</p> <p>» The <code>ssl_access_log</code> is used to track connections and for administrators to monitor the performance and security of the server and troubleshoot issues. This log can include:</p> <ul style="list-style-type: none"> » Date and time of the connection » IP address of the client making the connection » Certificate information presented by the client » Results of the SSL handshake process » Encryption algorithm and key size used for the connection » SSL protocol version used <p>» The <code>ssl_request_log</code> provides a more detailed record of</p>	

- » Results of processing the request
- » The **ssl_error_log** tracks errors related to SSL encryption. It records information about any issues that occur during the SSL handshake process or that prevent the establishment of a secure SSL connection.

Application Logs – MySQL

» Description

MySQL is a widely used open-source relational database management system. It also includes a comprehensive logging system, which by default involves writing to various log files, including the error log, the slow query log, and the general query log. MySQL allows configuration of the level of detail recorded in logs, although the default is basic logging and it may write to unexpected locations.

» Location

Typically, logging and log locations are defined by the configuration file at:

```
/etc/mysql/my.cnf
```

firewall is a common firewall management tool that is the default on RHEL/CentOS7 or later versions. It controls network traffic based on the rules specified in the firewall configuration, which can be changed dynamically without interrupting the network connections. The firewall also logs any events that occur if the “log” option is used.

» Location

```
/var/log/firewalld
```

» Interpretation

- » firewall does not log dropped packets by default; the “log” option is required
- » The default configuration is minimal logging
- » Logs can be very verbose, and therefore difficult to read
- » In mixed-mode, firewalld writes info-level log messages to syslog
- » Debug messages are written to the firewall log file by default, but this can be configured
- » Info messages also go to stdout and stderr

• Check the configuration file to locate logs and determine current configuration

- Log files are plaintext, so easily readable with a text editor or using the `cat` command

Application Logs – PostgreSQL

➤ Description

PostgreSQL is an open-source relational database management system that provides advanced features like transaction management, user authentication, and data integrity. PostgreSQL has a comprehensive logging system and can be configured to ship logs to the standard system log, as well as configure the level of detail that is recorded. However, the default logging options are often insufficient for security purposes and require deliberate effort to turn on.

➤ Location

Typically, logging and the log location is defined by the configuration file at:

```
# /etc/audit/audit.conf
```

Log location:

```
# /var/log/audit/audit.log
```

➤ Interpretation

- auditd is not installed by default on Debian-family distros
- Even where auditd is installed, the default settings are insufficient to provide much assistance to an investigation

- » `/var/lib/postgresql/data/postgresql.conf`
- » Interpretation
 - » Check the configuration file to locate logs and determine current configuration
 - » Log files are plaintext, so easily readable with a text editor or using the *cat* command
- » Search audit logs using the *ausearch* tool
- » Both *auseport* and *ausearch* query the live log on the running system by default. However, this can be overridden by specifying an input file with the *-if* argument

Computer Name

> Description

The computer name, or hostname, is stored in a plaintext file in the `/etc` folder.

> Location

- » `/etc/hostname`

> Interpretation

- » This file can be read with any text editor or via the `cat` command
- » Live response commands to identify the hostname:
 - » `hostname`
 - » `hostname -fqdn`
 - » `hostname -long`

Hosts File

> Description

The hosts file is used to map IP addresses with hostnames and takes precedence over a DNS lookup.

> Location

- » `/etc/hosts`

> Interpretation

- » This file can be read with any text editor or via the `cat` command
- » Check the hosts file for anything that may have been added/changed by an attacker, or if it provides any clues or visibility into the environment

Timezone

> Description

Only the device's current timezone setting is stored in these files.

> Location

- » `/etc/timezone`
- » `/etc/localtime` (symbolic link)

> Interpretation

- » `timezone` can be read with any text editor or via the `cat` command.
- » `localtime` can be read using the `ls -l` command.

Distro Information

> Description

The Linux distro name and version number are often recorded in multiple ways and formats, depending on the distro.

> Location

- » `/etc/os-release`

On Ubuntu (and Debian-based distros):

- » `/etc/lsb-release`

On RHEL/CentOS-based distros:

- » For RHEL systems: `/etc/redhat-release`
- » For Fedora systems: `/etc/fedora-release`
- » For CentOS systems: `/etc/centos-release`
- » For Rocky Linux systems: `/etc/rocky-release`
- » For Amazon Linux systems: `/etc/system-release`
- » For Oracle Linux systems: `/etc/oracle-release` and `/etc/redhat-release`
- » For SLES systems: `/etc/SuSE-release` (older) or `/etc/SUSE-brand` (newer)

> Interpretation

- » This file can be read with any text editor or via the `cat` command.
- » The key fields in `os-release` are:
 - » **NAME:** The distro name, which is defined by the creator/release organization
 - » **VERSION:** Version number and name, if appropriate
 - » **ID:** Also the distro name, but often in an abbreviated form
 - » **ID_LIKE:** The distro family
 - » **PRETTY_NAME:** A human-readable string describing the distro and version number
 - » **VERSION_ID:** A shortened representation of the version string
- » The key fields in `lsb-release` are:
 - » **DISTRIB_ID:** The distro name
 - » **DISTRIB_RELEASE:** The release or version number of the distro
 - » **DISTRIB_CODENAME:** The codename associated with the distro release
 - » **DISTRIB_DESCRIPTION:** A human-readable description of the distro, combining the name, version, and codename
- » RHEL/CentOS-based distro files contain a string with the distro name and version number
- » Live response commands to identify the distro information:
 - » `lsb_release -a`
 - » Note: This is not available on default Fedora or RHEL/CentOS 7+ distros.
 - » `hostnamectl`
 - » On remote systems (connecting via SSH): `hostnamectl -H [user]@[IP address]`
 - » `cat /proc/version`

Partition Information

> Description

Information about partitions currently present on the system's block devices, such as hard drives and flash drives, is available under the `/proc` virtual filesystem. This only exists while the system is live.

> Location

- » `/proc/partitions`

> Interpretation

- » Plaintext file; view using the `cat` command.

Mount Points

> Description

Information about the filesystems currently mounted in the live system, including their device, mount point, filesystem type, mount options, and other details, is present in a virtual file. This only exists while the system is live.

> Location

- » `/proc/mounts`

> Interpretation

- » Plaintext file; view using the `cat` command.
- » Does not provide a historical record for previously mounted filesystems.

